

OpenACC 2.5 Validation Testsuite targeting multiple architectures

Kyle Friedline, Sunita Chandrasekaran, University of Delaware, USA

Graham Lopez, Oscar Hernandez, Oak Ridge National Lab, USA

June 22, 2017

P3MA Workshop @ ISC

<https://github.com/OpenACCUserGroup/OpenACCV-V>

OpenACC- a directive-based parallel programming model

Add Simple Compiler Directive

```
main()
{
    <serial code>
    #pragma acc kernels
    {
        <parallel code>
    }
}
```



Motivation

- To help developers find and fix compiler bugs
- To assist users to know where the compilers stand in terms of OpenACC functionality support
- To assist users with fundamental understanding of spec features
- To discern ambiguities in the OpenACC specifications
- To discuss misinterpretations of feature definitions

Top HPC Apps adopting OpenACC

ANSYS Fluent
Gaussian VASP

3 of Top 10 Apps

GTC
XGC
ACME
FLASH
LSDalton

5 ORNL CAAR Codes

COSMO
ELEPHANT
RAMSES
ICON
ORB5

5 CSCS Codes

OpenACC

- Execution:
 - The `parallel` directive for explicit parallelization mapping
 - The `kernels` directive for implicit parallelization mapping
 - The `loop` construct for parallelization specification
- Memory:
 - The `data` directive for scoped, explicit-lifetime data elements
 - The `enter/exit` data directives for executable data transfer

OpenACC Execution Model

- **Parallel**

- Best when given explicit parallelization dimensions.

- **Kernels**

- Best when compiler is allowed to implement its determined best parallelization dimensions
- Potential for multiple compute kernels allowing multiple, varying-size compute-resource allocations.

- **Loop**

- Describes parallelization across for-loops with the addition of `gang`, `worker`, or `vector` clauses.
- Describes data dependencies both in code structure and in loop iterations through the `async`, `independent`, and `seq` clauses.

OpenACC Memory Model

- Enter/Exit data directives

- Executes memory allocation, deallocation, and transfer
- Allows for cross-scope data lifetimes
- Requires more vigilance on the part of the programmer to manage data lifetimes, especially with dynamic execution patterns

- Data region

- Gives data a specified beginning and ending on device
- Allows compiler to predetermine data presence via scopes and call trees.
- Can lead to data remaining allocated on device longer than necessary to complete scope

- Data management clauses

- Private
- Firstprivate
- Reduction
- Cache
- Deviceptr
- Present

Multi-Paradigm Accelerating Model

- Multiple Accelerator Support
 - `acc_get_num_devices`
 - `acc_set_device_type`
 - `acc_set_device_num`

Compilers supporting OpenACC

- PGI
 - NVIDIA GPU, X86 CPU, X86 Xeon Phi, POWER
- GNU
 - NVIDIA GPU
- Cray
 - NVIDIA GPU
- Sunway OpenACC*
- Omni (PEZY-SC) from University of Tsukuba
- OpenARC from ORNL, OpenUH from UH and SBU

<https://www.openacc.org/tools>

The Validation Test Suite

- The suite of 177 tests includes 86 Fortran tests and 91 C tests covering:
 - All of the parallel construct and its clauses, but missing some fringe functionality and clause combinations
 - Most of kernels construct and its clauses, but missing all fringe functionality and clause combinations
 - Data directives (data, enter data, exit data) with primary clauses tested, but missing all fringe functionality and most clause combinations
- Tests build upon each other; initial tests test core functionality, while later tests built upon previously tested functionality

Test Suite Structure

- Primary test types
 - Unit tests
 - Incremental Integration Tests
 - Use Cases
- Test Design Issues
 - Backwards-compatibility
 - Compute complexity
 - Platform-independent test compatibility

Test Design

```
double * a = (double *)malloc(n * sizeof(double));
double * b = (double *)malloc(n * sizeof(double));
double * a_copy = (double *)malloc(n * sizeof(double));
int * devtest = (int *)malloc(sizeof(int));
int err = 0;
devtest[0] = 1;
#pragma acc enter data copyin(devtest[0:1])
#pragma acc parallel present(devtest[0:1])
{
    devtest[0] = 0;
}
for (int x = 0; x < n; ++x){
    a[x] = rand();
    a_copy[x] = a[x];
}
if (devtest[0] == 1 && run_probabilistic == 1){
    #pragma acc enter data create(a[0:n])
    #pragma acc exit data copyout(a[0:n])
    err = 1; \\Failing
    for (int x = 0; x < n; ++x){
        if (fabs(a[x] - a_copy[x]) > PRECISION){
            err = 0;
        }
    }
}
for (int x = 0; x < n; ++x){
    a[x] = rand();
    b[x] = 0.0;
}
#pragma acc enter data copyin(a[0:n]) create(b[0:n])
#pragma acc parallel present(a[0:n], b[0:n])
{
    for (int x = 0; x < n; ++x){
        b[x] = a[x];
    }
}
#pragma acc exit data copyout(b[0:n]) delete(a[0:n])
for (int x = 0; x < n; ++x){
    if (fabs(a[x] - b[x]) > PRECISION){
        err += 1;
        break;
    }
}
```

- Design issues addressed
 - Data clauses are called as 2.5 specifies
 - Compute complexity is kept minimal
 - Internal testing determines memory type for some platform dependent testing

Contributions of the Test Suite

- Check for correctness of implementations of OpenACC features
- Reported bugs to both PGI and GCC
- Check and report ambiguities or misinterpretation in the OpenACC specification

Open-sourcing V&V Suite

- Validation and Verification Suite will soon be made available to everyone
 - Target deadline late August 2017
- <https://github.com/OpenACCUUserGroup/OpenACCV-V>
- Feedback on the suite is welcome
- Contributions to the suite welcome
- BSD 3-Clause license

Tested Platforms

	K20	K80	P100	Ivy Bridge	Bulldoz er	Power8 +	Knights Landing
PGI	V16.10	V16.10*	V17.1	V16.10** /17.3	V16.10	V17.1	V17.1***
GNU	V6.3-20 170303	V6.0.0-20 160415		V6.0.0-20 160415	V6.3-201 70303		

*K80 was run against many versions of PGI which will be seen in a following slide.

**Version 16.10 reported in preference of 17.3 since 16.10 was released as the community edition. Freely available at no cost for download

***Knights Landing is not officially supported; testing was performed by using `-ta=haswell` flag.

Results from PGI and GNU compilers

Architecture	PGI Pass Rate	GNU Pass Rate
K20	175/177	112/177
K80	175/177	113/177
Ivy Bridge	171/177	154/177
Bulldozer	172/177	157/177

Cross-Platform Performance of PGI

Architecture	Pass Rate
K20	175/177
K80	175/177
P100	175/177
Ivy Bridge	171/177
Bulldozer	172/177
Power8+	165/177
Knights Landing	167/177

PGI Improvement Over Versions

Compiler Version	Fortran Pass Rate	C Pass Rate	Fortran % Passed	C % Passed
14.10	60/86	67/91	69.8	73.6
15.1	64/86	80/91	74.4	87.9
15.5	65/86	80/91	95.6	87.9
15.10	68/86	84/91	79.1	92.3
16.1	69/86	84/91	80.2	92.3
16.4	82/86	84/91	95.3	92.3
16.7	85/86	90/91	98.8	98.9
16.10	85/86	90/91	98.8	98.9
17.1	85/86	90/91	98.8	98.9
17.3	85/86	90/91	98.8	98.9

These results are from running against K80

Future Work

- Near future work – open source V&V suite
- Improve documentation
 - Forward and backward referencing
- Addition of use-cases
- Build a more comprehensive suite (cross, orphan test cases)
- Keep up with the growing feature set of the specification
- Brainstorming building an infrastructure common to both OpenMP and OpenACC V&V